

svm-regression-me

December 4, 2025

SVM_REG

```
[2]: ### Importation des bibliothèques
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
from sklearn.svm import SVR
import joblib
import warnings
warnings.filterwarnings('ignore')

print("=="*70)
print("SVM REGRESSION - CAR PRICE PREDICTION")
print("=="*70)

### Chargement des données
df = pd.read_csv("Car_Price_Prediction.csv")
print(f"\nDataset loaded: {df.shape[0]} rows, {df.shape[1]} columns")

### Exploration des données
print("\n" + "=="*70)
print("DATA EXPLORATION")
print("=="*70)
print("\nFirst 5 rows:")
print(df.head())

print("\nDataset Info:")
print(df.info())

print("\nStatistical Summary:")
print(df.describe())

print("\nMissing Values:")
print(df.isna().sum())
```

```

### Nettoyage des données (outliers)
print("\n" + "="*70)
print("DATA CLEANING - OUTLIER DETECTION")
print("="*70)

num_cols = df.select_dtypes(include=["int", "float"]).columns

for col in num_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = (df[col] < lower_bound) | (df[col] > upper_bound)
    print(f"{col}: {outliers.sum()} outliers detected")
    df[col] = df[col].clip(lower=lower_bound, upper=upper_bound)

### Encodage des variables catégorielles
print("\n" + "="*70)
print("CATEGORICAL ENCODING")
print("="*70)

# OneHotEncoder pour Make, Model, Fuel Type
cat_multi_cols = ["Make", "Model", "Fuel Type"]
print(f"\nApplying OneHotEncoder to: {cat_multi_cols}")

OHE = OneHotEncoder()
col_ohe = OHE.fit_transform(df[cat_multi_cols])
col_ohe = pd.DataFrame(
    col_ohe.toarray(),
    columns=OHE.get_feature_names_out(cat_multi_cols),
    dtype='int'
)
print(f"One-hot encoded features: {col_ohe.shape[1]}")

# LabelEncoder pour Transmission
le = LabelEncoder()
df['Transmission'] = le.fit_transform(df['Transmission'])
print(f"\nTransmission encoding: {dict(enumerate(le.classes_))}")

# Reconstruire le dataframe
df = df.drop(cat_multi_cols, axis=1)
df = pd.concat([col_ohe, df], axis=1)

print(f"\nFinal dataset shape: {df.shape}")
print(f"Total features (including target): {len(df.columns)}")

```

```

print(f"\nColumn order:")
for i, col in enumerate(df.columns):
    print(f"  [{i}] {col}")

### Préparation des données pour l'entraînement
print("\n" + "="*70)
print("TRAIN-TEST SPLIT")
print("="*70)

X = df.drop("Price", axis=1)
y = df["Price"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

print(f"Training set: {X_train.shape}")
print(f"Test set: {X_test.shape}")
print(f"Features: {X_train.shape[1]}")

### Normalisation des données
print("\n" + "="*70)
print("FEATURE SCALING (StandardScaler)")
print("="*70)

# For SVM, we need to scale the numerical features
# Features to scale: Year (index 13), Engine Size (index 14), Mileage (index 15)
# One-hot features (0-12) and Transmission (16) should NOT be scaled

cols_to_scale_indices = [13, 14, 15] # Year, Engine Size, Mileage
cols_to_scale_names = ['Year', 'Engine Size', 'Mileage']

print(f"Scaling features: {cols_to_scale_names}")
print(f"Indices: {cols_to_scale_indices}")

scaler = StandardScaler()

# Extract only the columns to scale for training
X_train_to_scale = X_train.iloc[:, cols_to_scale_indices]
X_test_to_scale = X_test.iloc[:, cols_to_scale_indices]

# Fit and transform
scaler.fit(X_train_to_scale)
X_train_scaled = scaler.transform(X_train_to_scale)
X_test_scaled = scaler.transform(X_test_to_scale)

# Replace scaled columns back into the original dataframes
X_train_final = X_train.copy()

```

```

X_test_final = X_test.copy()

X_train_final.iloc[:, cols_to_scale_indices] = X_train_scaled
X_test_final.iloc[:, cols_to_scale_indices] = X_test_scaled

print(f"\nScaler fitted on training data")
print(f"Scaler mean: {scaler.mean_}")
print(f"Scaler std: {scaler.scale_}")

### Entraînement du modèle SVM de base
print("\n" + "="*70)
print("MODEL TRAINING - BASIC SVM REGRESSION")
print("="*70)

# SVM avec kernel RBF (Radial Basis Function)
svm_basic = SVR(kernel='rbf', C=100, gamma='scale', epsilon=0.1)

print("Training SVM model...")
svm_basic.fit(X_train_final, y_train)
print(" Model trained successfully!")

y_pred_basic = svm_basic.predict(X_test_final)

### Évaluation du modèle de base
def evaluate_model(y_true, y_pred, model_name):
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(y_true, y_pred)
    r2 = r2_score(y_true, y_pred)

    print(f"\n{model_name}:")
    print(f"  MSE: {mse:.2f}")
    print(f"  RMSE: {rmse:.2f}")
    print(f"  MAE: {mae:.2f}")
    print(f"  R²: {r2:.4f}")

    return {'MSE': mse, 'RMSE': rmse, 'MAE': mae, 'R2': r2}

metrics_basic = evaluate_model(y_test, y_pred_basic, "SVM Regression - Basic Model")

### Optimisation des hyperparamètres avec GridSearchCV
print("\n" + "="*70)
print("HYPERPARAMETER OPTIMIZATION (GridSearchCV)")
print("="*70)

param_grid = {

```

```

'C': [10, 50, 100, 200],
'gamma': ['scale', 'auto', 0.001, 0.01, 0.1],
'epsilon': [0.01, 0.1, 0.2],
'kernel': ['rbf']
}

print(f"Parameter grid: {param_grid}")
print("\nSearching for best parameters... (this may take a few minutes)")

svm_model = SVR()
grid_search = GridSearchCV(
    svm_model,
    param_grid,
    cv=5,
    scoring='neg_mean_squared_error',
    n_jobs=-1,
    verbose=1
)

grid_search.fit(X_train_final, y_train)

print("\n Grid search completed!")
print(f"Best parameters: {grid_search.best_params_}")
print(f"Best cross-validation score: {-grid_search.best_score_:.2f} (MSE)")

### Modèle optimisé
print("\n" + "="*70)
print("OPTIMIZED MODEL EVALUATION")
print("="*70)

best_svm = grid_search.best_estimator_
y_pred_optimized = best_svm.predict(X_test_final)

metrics_optimized = evaluate_model(y_test, y_pred_optimized, "SVM Regression - Optimized Model")

### Comparaison des modèles
print("\n" + "="*70)
print("MODEL COMPARISON")
print("="*70)

comparison = pd.DataFrame({
    'Metric': ['MSE', 'RMSE', 'MAE', 'R²'],
    'Basic Model': [
        f"{metrics_basic['MSE']:.2f}",
        f"{metrics_basic['RMSE']:.2f}",
        f"{metrics_basic['MAE']:.2f}",

```

```

        f"{{metrics_basic['R2']:.4f}}"
    ],
'Optimized Model': [
    f"{{metrics_optimized['MSE']:.2f}}",
    f"{{metrics_optimized['RMSE']:.2f}}",
    f"{{metrics_optimized['MAE']:.2f}}",
    f"{{metrics_optimized['R2']:.4f}}"
]
})
print(comparison.to_string(index=False))

# Calculate improvement
r2_improvement = ((metrics_optimized['R2'] - metrics_basic['R2']) / 
    metrics_basic['R2']) * 100
print(f"\nR2 Improvement: {r2_improvement:+.2f}%")

### Visualisation des prédictions
print("\n" + "="*70)
print("VISUALIZATION")
print("="*70)

fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# Plot 1: Actual vs Predicted (Optimized Model)
axes[0].scatter(y_test, y_pred_optimized, alpha=0.5, edgecolors='k')
axes[0].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', 
    lw=2)
axes[0].set_xlabel('Actual Price ($)')
axes[0].set_ylabel('Predicted Price ($)')
axes[0].set_title('SVM Regression: Actual vs Predicted Prices')
axes[0].grid(True, alpha=0.3)

# Plot 2: Residuals
residuals = y_test - y_pred_optimized
axes[1].scatter(y_pred_optimized, residuals, alpha=0.5, edgecolors='k')
axes[1].axhline(y=0, color='r', linestyle='--', lw=2)
axes[1].set_xlabel('Predicted Price ($)')
axes[1].set_ylabel('Residuals ($)')
axes[1].set_title('Residual Plot')
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('svm_regression_results.png', dpi=300, bbox_inches='tight')
print("\n Visualization saved as 'svm_regression_results.png'")
plt.show()

### Sauvegarde du modèle optimisé

```

```

print("\n" + "="*70)
print("MODEL SAVING")
print("="*70)

# Save the optimized model
joblib.dump(best_svm, "SVM_Regression.pkl")
print(" Model saved: SVM_Regression.pkl")

# Save the scaler (very important for SVM!)
joblib.dump(scaler, "Scaler_SVM_Regression.pkl")
print(" Scaler saved: Scaler_SVM_Regression.pkl")

print("\n" + "="*70)
print("IMPORTANT NOTES FOR DJANGO INTEGRATION")
print("="*70)
print"""

1. Feature Order (17 features total):
    [0-4] Make (one-hot): Audi, BMW, Ford, Honda, Toyota
    [5-9] Model (one-hot): Model A, B, C, D, E
    [10-12] Fuel Type (one-hot): Diesel, Electric, Petrol
    [13] Year (SCALED)
    [14] Engine Size (SCALED)
    [15] Mileage (SCALED)
    [16] Transmission (NOT SCALED, 0=Automatic, 1=Manual)

2. The scaler is fitted ONLY on indices [13, 14, 15]

3. Django must apply the same transformation:
    - One-hot encode categorical variables
    - Scale only Year, Engine Size, Mileage
    - Keep Transmission as 0 or 1

4. Files needed in Django models_ai/ folder:
    - SVM_Regression.pkl
    - Scaler_SVM_Regression.pkl
""")

print("\n" + "="*70)
print("SAMPLE PREDICTION TEST")
print("="*70)

# Test with a sample
sample_features = X_test_final.iloc[0:1]
sample_actual = y_test.iloc[0]
sample_pred = best_svm.predict(sample_features)[0]

print(f"\nSample Test:")

```

```

print(f" Actual Price:    ${sample_actual:,.2f}")
print(f" Predicted Price: ${sample_pred:,.2f}")
print(f" Difference:      ${abs(sample_actual - sample_pred):,.2f}")
print(f" Error:           {abs(sample_actual - sample_pred) / sample_actual * 100:.2f}%")


print("\n" + "="*70)
print(" SVM REGRESSION TRAINING COMPLETED SUCCESSFULLY!")
print("="*70)

```

```
=====
SVM REGRESSION - CAR PRICE PREDICTION
=====
```

Dataset loaded: 1000 rows, 8 columns

```
=====
DATA EXPLORATION
=====
```

First 5 rows:

	Make	Model	Year	Engine Size	Mileage	Fuel Type	Transmission	\
0	Honda	Model B	2015	3.9	74176	Petrol	Manual	
1	Ford	Model C	2014	1.7	94799	Electric	Automatic	
2	BMW	Model B	2006	4.1	98385	Electric	Manual	
3	Honda	Model B	2015	2.6	88919	Electric	Automatic	
4	Honda	Model C	2004	3.4	138482	Petrol	Automatic	

	Price
0	30246.207931
1	22785.747684
2	25760.290347
3	25638.003491
4	21021.386657

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Make	1000 non-null	object
1	Model	1000 non-null	object
2	Year	1000 non-null	int64
3	Engine Size	1000 non-null	float64
4	Mileage	1000 non-null	int64
5	Fuel Type	1000 non-null	object

```
6    Transmission    1000 non-null    object
7    Price           1000 non-null    float64
dtypes: float64(2), int64(2), object(4)
memory usage: 62.6+ KB
None
```

Statistical Summary:

	Year	Engine Size	Mileage	Price
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	2010.688000	2.798300	97192.48700	25136.615530
std	6.288577	1.024137	59447.31576	5181.401368
min	2000.000000	1.000000	56.000000	6704.953524
25%	2005.000000	1.900000	44768.75000	21587.878370
50%	2011.000000	2.800000	94411.50000	25189.325247
75%	2016.000000	3.700000	148977.75000	28806.368974
max	2021.000000	4.500000	199867.00000	41780.504635

Missing Values:

```
Make          0
Model         0
Year          0
Engine Size   0
Mileage        0
Fuel Type      0
Transmission   0
Price          0
dtype: int64
```

DATA CLEANING - OUTLIER DETECTION

```
Year: 0 outliers detected
Engine Size: 0 outliers detected
Mileage: 0 outliers detected
Price: 3 outliers detected
```

CATEGORICAL ENCODING

```
Applying OneHotEncoder to: ['Make', 'Model', 'Fuel Type']
One-hot encoded features: 13
```

```
Transmission encoding: {0: 'Automatic', 1: 'Manual'}
```

```
Final dataset shape: (1000, 18)
Total features (including target): 18
```

```
Column order:  
[0] Make_Audi  
[1] Make_BMW  
[2] Make_Ford  
[3] Make_Honda  
[4] Make_Toyota  
[5] Model_Model A  
[6] Model_Model B  
[7] Model_Model C  
[8] Model_Model D  
[9] Model_Model E  
[10] Fuel Type_Diesel  
[11] Fuel Type_Electric  
[12] Fuel Type_Petrol  
[13] Year  
[14] Engine Size  
[15] Mileage  
[16] Transmission  
[17] Price
```

```
=====  
TRAIN-TEST SPLIT  
=====
```

```
Training set: (800, 17)  
Test set: (200, 17)  
Features: 17
```

```
=====  
FEATURE SCALING (StandardScaler)  
=====
```

```
Scaling features: ['Year', 'Engine Size', 'Mileage']  
Indices: [13, 14, 15]
```

```
Scaler fitted on training data  
Scaler mean: [2.01084000e+03 2.80800000e+00 9.83993875e+04]  
Scaler std: [6.21987942e+00 1.03310745e+00 5.94083977e+04]
```

```
=====  
MODEL TRAINING - BASIC SVM REGRESSION  
=====
```

```
Training SVM model...  
Model trained successfully!
```

```
SVM Regression - Basic Model:  
MSE: 10,366,887.92  
RMSE: 3,219.77  
MAE: 2,645.12  
R2: 0.6212
```

```
=====
HYPERPARAMETER OPTIMIZATION (GridSearchCV)
=====
Parameter grid: {'C': [10, 50, 100, 200], 'gamma': ['scale', 'auto', 0.001, 0.01, 0.1], 'epsilon': [0.01, 0.1, 0.2], 'kernel': ['rbf']}
Searching for best parameters... (this may take a few minutes)
Fitting 5 folds for each of 60 candidates, totalling 300 fits

    Grid search completed!
Best parameters: {'C': 200, 'epsilon': 0.01, 'gamma': 'auto', 'kernel': 'rbf'}
Best cross-validation score: 6,238,896.63 (MSE)

=====
OPTIMIZED MODEL EVALUATION
=====

SVM Regression - Optimized Model:
    MSE: 6,018,715.19
    RMSE: 2,453.31
    MAE: 2,003.30
    R2: 0.7801

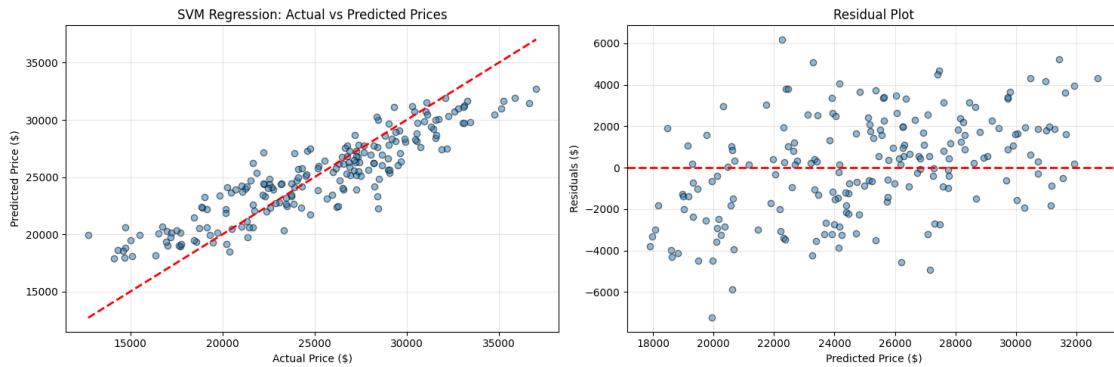
=====
MODEL COMPARISON
=====

Metric      Basic Model Optimized Model
    MSE 10,366,887.92      6,018,715.19
    RMSE      3,219.77      2,453.31
    MAE       2,645.12      2,003.30
    R2        0.6212      0.7801

R2 Improvement: +25.58%

=====
VISUALIZATION
=====

    Visualization saved as 'svm_regression_results.png'
```



MODEL SAVING

```
Model saved: SVM_Regression.pkl
Scaler saved: Scaler_SVM_Regression.pkl
```

IMPORTANT NOTES FOR DJANGO INTEGRATION

1. Feature Order (17 features total):
 - [0-4] Make (one-hot): Audi, BMW, Ford, Honda, Toyota
 - [5-9] Model (one-hot): Model A, B, C, D, E
 - [10-12] Fuel Type (one-hot): Diesel, Electric, Petrol
 - [13] Year (SCALED)
 - [14] Engine Size (SCALED)
 - [15] Mileage (SCALED)
 - [16] Transmission (NOT SCALED, 0=Automatic, 1=Manual)
2. The scaler is fitted ONLY on indices [13, 14, 15]
3. Django must apply the same transformation:
 - One-hot encode categorical variables
 - Scale only Year, Engine Size, Mileage
 - Keep Transmission as 0 or 1
4. Files needed in Django models_ai/ folder:
 - SVM_Regression.pkl
 - Scaler_SVM_Regression.pkl

SAMPLE PREDICTION TEST

=====

Sample Test:

Actual Price: \$27,902.29

Predicted Price: \$25,636.10

Difference: \$2,266.19

Error: 8.12%

=====

SVM REGRESSION TRAINING COMPLETED SUCCESSFULLY!

=====