

# XGBoost Classification for Heart Disease Prediction (Without Scaling)

This notebook builds an XGBoost classifier to predict heart disease using the synthetic dataset. We will reduce to the top 7 features based on importance, without scaling.

## Step 1: Import Libraries

We import necessary libraries for data handling, preprocessing, modeling, and evaluation.

```
In [10]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import xgboost as xgb
import pickle
```

## Step 2: Load the Dataset

Load the CSV file into a pandas DataFrame and display basic information.

```
In [11]: df = pd.read_csv('synthetic_heart_disease_dataset.csv')
print("Dataset shape:", df.shape)
print("Columns:", df.columns.tolist())
print("First 5 rows:")
print(df.head())
print("\nData types:")
print(df.dtypes)
print("\nMissing values:")
print(df.isnull().sum())
```

Dataset shape: (50000, 21)  
Columns: ['Age', 'Gender', 'Weight', 'Height', 'BMI', 'Smoking', 'Alcohol\_Intake', 'Physical\_Activity', 'Diet', 'Stress\_Level', 'Hypertension', 'Diabetes', 'Hyperlipidemia', 'Family\_History', 'Previous\_Heart\_Attack', 'Systolic\_BP', 'Diastolic\_BP', 'Heart\_Rate', 'Blood\_Sugar\_Fasting', 'Cholesterol\_Total', 'Heart\_Disease']

First 5 rows:

```
Age  Gender  Weight  Height  BMI  Smoking Alcohol_Intake \
0   48     Male    78      157  26.4  Never      NaN
1   35   Female   73      163  33.0  Never      Low
2   79   Female   88      152  32.3  Never      NaN
3   75     Male   106      171  37.4  Never  Moderate
4   34   Female   65      191  18.5  Current      NaN
```

```
Physical_Activity      Diet Stress_Level ... Diabetes Hyperlipidemia \
0       Sedentary  Healthy      Medium ... 0          1
1       Active    Average      High   ... 0          1
2     Moderate  Average      Medium ... 0          0
3     Moderate  Average      Low    ... 0          1
4     Sedentary  Healthy      Low    ... 1          0
```

```
Family_History  Previous_Heart_Attack  Systolic_BP  Diastolic_BP \
0            1                  0        104          99
1            1                  0        111          72
2            1                  0        116         102
3            0                  0        171          92
4            0                  0        164          67
```

```
Heart_Rate  Blood_Sugar_Fasting  Cholesterol_Total  Heart_Disease
0        71                  165            200          0
1        60                  145            206          0
2        78                  148            208          0
3       109                  105            290          1
4       108                  116            220          1
```

[5 rows x 21 columns]

Data types:

Age	int64
Gender	object
Weight	int64
Height	int64
BMI	float64
Smoking	object
Alcohol_Intake	object
Physical_Activity	object
Diet	object
Stress_Level	object
Hypertension	int64
Diabetes	int64
Hyperlipidemia	int64
Family_History	int64
Previous_Heart_Attack	int64
Systolic_BP	int64
Diastolic_BP	int64
Heart_Rate	int64
Blood_Sugar_Fasting	int64

```
Cholesterol_Total           int64
Heart_Disease               int64
dtype: object

Missing values:
Age                          0
Gender                       0
Weight                        0
Height                        0
BMI                           0
Smoking                      0
Alcohol_Intake              20109
Physical_Activity            0
Diet                          0
Stress_Level                 0
Hypertension                  0
Diabetes                      0
Hyperlipidemia                0
Family_History                0
Previous_Heart_Attack        0
Systolic_BP                   0
Diastolic_BP                  0
Heart_Rate                     0
Blood_Sugar_Fasting           0
Cholesterol_Total             0
Heart_Disease                  0
dtype: int64
```

## Step 3: Preprocess the Data

Handle missing values by filling with mode for categorical, encode categorical variables to numerical.

```
In [12]: # Handle missing values
df['Alcohol_Intake'].fillna(df['Alcohol_Intake'].mode()[0], inplace=True)

# Encode categorical variables
categorical_cols = ['Gender', 'Smoking', 'Alcohol_Intake', 'Physical_Activity', 'Diet']
le = LabelEncoder()
for col in categorical_cols:
    df[col] = le.fit_transform(df[col])

print("After preprocessing:")
print(df.head())
```

After preprocessing:

```
   Age  Gender  Weight  Height  BMI  Smoking  Alcohol_Intake \
0    48       1      78     157  26.4        2             1
1    35       0      73     163  33.0        2             1
2    79       0      88     152  32.3        2             1
3    75       1     106     171  37.4        2             2
4    34       0      65     191  18.5        0             1

   Physical_Activity  Diet  Stress_Level  ...  Diabetes  Hyperlipidemia \
0                  2     1                 2  ...          0             1
1                  0     0                 0  ...          0             1
2                  1     0                 2  ...          0             0
3                  1     0                 1  ...          0             1
4                  2     1                 1  ...          1             0

   Family_History  Previous_Heart_Attack  Systolic_BP  Diastolic_BP \
0            1                   0           104            99
1            1                   0           111            72
2            1                   0           116           102
3            0                   0           171            92
4            0                   0           164            67

   Heart_Rate  Blood_Sugar_Fasting  Cholesterol_Total  Heart_Disease
0       71                  165            200              0
1       60                  145            206              0
2       78                  148            208              0
3      109                  105            290              1
4      108                  116            220              1
```

[5 rows x 21 columns]

C:\Users\yahya\AppData\Local\Temp\ipykernel\_24524\2640066891.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Alcohol_Intake'].fillna(df['Alcohol_Intake'].mode()[0], inplace=True)
```

## Step 4: Feature Selection Function

Define a function to train an initial XGBoost model and select the top 7 features based on gain importance.

```
In [21]: def feature_importance_function(X, y, n_features=8):
    """
    Computes feature importance using XGBoost and returns the top n features.
    """
    model = xgb.XGBClassifier(objective='binary:logistic', n_estimators=100, random
```

```

importance = model.get_booster().get_score(importance_type='gain')
sorted_features = sorted(importance.items(), key=lambda x: x[1], reverse=True)
print("Top", n_features, "features by importance:")
for i, (feature, score) in enumerate(sorted_features[:n_features]):
    print(f"{i+1}. {feature}: {score:.4f}")
top_features = [f[0] for f in sorted_features[:n_features]]
return top_features

X = df.drop('Heart_Disease', axis=1)
y = df['Heart_Disease']
top_features = feature_importance_function(X, y, 8)
X_selected = X[top_features]

```

Top 8 features by importance:

1. Hypertension: 334.2283
2. Cholesterol\_Total: 271.0656
3. Diabetes: 248.6891
4. Age: 230.3551
5. Previous\_Heart\_Attack: 126.4948
6. Systolic\_BP: 0.0023
7. Heart\_Rate: 0.0007
8. Diastolic\_BP: 0.0001

## Step 5: Split the Data

Split the selected features and target into training and testing sets.

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.2, r
print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)
```

Training set shape: (40000, 7)  
Testing set shape: (10000, 7)

## Step 6: Train the XGBoost Model

Initialize and train the XGBoost classifier on the training data.

```
In [15]: model = xgb.XGBClassifier(
    objective='binary:logistic',
    n_estimators=100,
    max_depth=6,
    learning_rate=0.1,
    random_state=42
)
model.fit(X_train, y_train)
print("Model trained.")
```

Model trained.

## Step 7: Evaluate the Model

Make predictions on the test set and evaluate performance.

```
In [16]: y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 1.0000

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	5365
1	1.00	1.00	1.00	4635
accuracy			1.00	10000
macro avg	1.00	1.00	1.00	10000
weighted avg	1.00	1.00	1.00	10000

Confusion Matrix:

```
[[5365  0]
 [ 0 4635]]
```

## Step 8: Save the Model

Save the trained model to a pickle file.

```
In [17]: with open('xgboost_without_scaler.pkl', 'wb') as f:
    pickle.dump(model, f)
print("Model saved to xgboost_without_scaler.pkl")
```

Model saved to xgboost\_without\_scaler.pkl